


Perturbed optimizers for learning



Q. Berthet
(Google DeepMind)




“Differentiable Almost Everything” workshop

ICML 2023, Honolulu, Hawaii



End-to-end differentiable models


ML training: minimizing a loss w.r.t. weights of a model



- Composition of complex, explicit, analytical operations with **weights**.
- Loss minimization with first-order, **gradient-based** methods.
- Modern, large-scale models and datasets: **automatic differentiation**.


End-to-end differentiable models

ML training: minimizing a loss w.r.t. weights of a model



- Composition of complex, explicit, analytical operations with **weights**.
- Loss minimization with first-order, **gradient-based** methods.
- Modern, large-scale models and datasets: **automatic differentiation**.


Discrete operations



- **Discrete** operations / algorithms at the heart of computer science.
- Powerful tool to deal with **structured** problems / data.

Structured inference and prediction


Discrete operators allow us to leverage **structure**.



- Use of an **implicit** and **discrete** function, challenge for **gradients**
- Challenge: moving embeddings towards “correct solution” continuously
- In **classification**, soft “arg” max solution, smooth approximation.

Structured inference and prediction



Discrete operators allow us to leverage **structure**.




- Use of an **implicit** and **discrete** function, challenge for **gradients**
- Challenge: moving embeddings towards “correct solution” continuously
- In **classification**, soft “arg” max solution, smooth approximation.

Structured inference and prediction

Discrete operators allow us to leverage **structure**.



Loss




Ground truth

- Use of an **implicit** and **discrete** function, challenge for **gradients**
- Challenge: moving embeddings towards “correct solution” continuously
- In **classification**, soft “arg” max solution, smooth approximation.

Structured inference and prediction


Discrete operators allow us to leverage **structure**.



- Use of an **implicit** and **discrete** function, challenge for **gradients**
- Challenge: moving embeddings towards “correct solution” continuously
- In **classification**, soft “arg” max solution, smooth approximation.


Structured inference and prediction

Discrete operators allow us to leverage **structure**.



- Use of an **implicit** and **discrete** function, challenge for **gradients**
- Challenge: moving embeddings towards “correct solution” continuously
- In **classification**, soft “arg” max solution, smooth approximation.


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / perturbed variational definition, with **closed form**

$$y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle, \quad [y^*(\theta)]_i = \delta_{i^*}(i).$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / perturbed variational definition, with **closed form**

$$y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle, \quad [y^*(\theta)]_i = \delta_{i^*}(i).$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / perturbed variational definition, with **closed form**

$$y^*(\theta) = \underset{y \in \mathcal{C}}{\operatorname{argmax}} \langle y, \theta \rangle, \quad [y^*(\theta)]_i = \delta_{i^*}(i).$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent **regularized** / perturbed variational definition, with **closed form**

$$y_\varepsilon^*(\theta) = \underset{y \in \mathcal{C}}{\operatorname{argmax}} \langle y, \theta \rangle - \varepsilon \Omega(y), \quad [y_\varepsilon^*(\theta)]_i = \frac{e^{\theta_i/\varepsilon}}{\sum_j e^{\theta_j/\varepsilon}}.$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / **perturbed** variational definition, with **closed form**

$$y_\varepsilon^*(\theta) = \mathbf{E}_{y \in \mathcal{C}}[\operatorname{argmax} \langle y, \theta + \varepsilon Z \rangle], \quad [y_\varepsilon^*(\theta)]_i = \frac{e^{\theta_i/\varepsilon}}{\sum_j e^{\theta_j/\varepsilon}}.$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / **perturbed** variational definition, with **closed form**

$$y_\varepsilon^*(\theta) = \underset{y \in \mathcal{C}}{\mathbf{E}}[\operatorname{argmax} \langle y, \theta + \varepsilon Z \rangle], \quad [y_\varepsilon^*(\theta)]_i = \frac{e^{\theta_i/\varepsilon}}{\sum_j e^{\theta_j/\varepsilon}}.$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / **perturbed** variational definition, with **closed form**

$$y_\varepsilon^*(\theta) = \underset{y \in \mathcal{C}}{\text{argmax}} \langle y, \theta + \varepsilon Z \rangle, \quad [y_\varepsilon^*(\theta)]_i = \frac{e^{\theta_i/\varepsilon}}{\sum_j e^{\theta_j/\varepsilon}}.$$


Softmax



- **Softmax** a smooth, explicit approximation of “hard” argmax over **simplex**.
- Equivalent regularized / **perturbed** variational definition, with **closed form**

$$y_\varepsilon^*(\theta) = \underset{y \in \mathcal{C}}{\mathbf{E}}[\operatorname{argmax}\langle y, \theta + \varepsilon Z \rangle], \quad [y_\varepsilon^*(\theta)]_i = \frac{e^{\theta_i/\varepsilon}}{\sum_j e^{\theta_j/\varepsilon}}.$$

Generalized softmax (?)




- Most discrete optimizers can be naturally written as

$$y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle$$

- How to generalize these methods, to have **differentiable** proxies?

Generalized softmax (?)




- Most discrete optimizers can be naturally written as

$$y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle$$

- How to generalize these methods, to have **differentiable** proxies?

Generalized softmax (?)




- Most discrete optimizers can be naturally written as

$$y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle$$

- How to generalize these methods, to have **differentiable** proxies?

Generalized softmax (?)



- Most discrete optimizers can be naturally written as

$$y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle$$

- How to generalize these methods, to have **differentiable** proxies?



Q. Berthet



M. Blondel



O.Teboul



M. Cuturi



J-P. Vert



F.Bach


- Learning with Differentiable Perturbed Optimizers

NeurIPS 2020

Perturbed maximizer

Discrete decisions: optimizers of linear program over \mathcal{C} , convex hull of $\mathcal{Y} \subseteq \mathbf{R}^k$

$$F(\theta) = \max_{y \in \mathcal{C}} \langle y, \theta \rangle, \quad \text{and} \quad y^*(\theta) = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta \rangle = \nabla_\theta F(\theta).$$




Perturbed maximizer: average of solutions for inputs with noise εZ

$$F_\varepsilon(\theta) = \mathbf{E}[\max_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle], \quad y_\varepsilon^*(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[\operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle] = \nabla_\theta F_\varepsilon(\theta).$$

Perturbed maximizer

Discrete decisions: optimizers of linear program over \mathcal{C} , convex hull of $\mathcal{Y} \subseteq \mathbb{R}^k$




Perturbed maximizer: average of solutions for inputs with noise εZ

$$F_\varepsilon(\theta) = \mathbf{E}[\max_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle], \quad y_\varepsilon^*(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[\operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle] = \nabla_\theta F_\varepsilon(\theta).$$

Perturbed maximizer

Discrete decisions: optimizers of linear program over \mathcal{C} , convex hull of $\mathcal{Y} \subseteq \mathbf{R}^k$




Perturbed maximizer: average of solutions for inputs with noise εZ

$$F_\varepsilon(\theta) = \mathbf{E}[\max_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle], \quad y_\varepsilon^*(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[\operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle] = \nabla_\theta F_\varepsilon(\theta).$$

Perturbed maximizer

Discrete decisions: optimizers of linear program over \mathcal{C} , convex hull of $\mathcal{Y} \subseteq \mathbb{R}^k$




Perturbed maximizer: average of solutions for inputs with noise εZ

$$F_\varepsilon(\theta) = \mathbf{E}[\max_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle], \quad y_\varepsilon^*(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[\operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle] = \nabla_\theta F_\varepsilon(\theta).$$

Perturbed maximizer

Discrete decisions: optimizers of linear program over \mathcal{C} , convex hull of $\mathcal{Y} \subseteq \mathbb{R}^k$



Perturbed maximizer: average of solutions for inputs with noise εZ

$$F_\varepsilon(\theta) = \mathbf{E}[\max_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle], \quad y_\varepsilon^*(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[\operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle] = \nabla_\theta F_\varepsilon(\theta).$$


Perturbed model

Model of optimal decision under uncertainty Luce (1959), McFadden et al. (1973)

$$Y = \operatorname{argmax}_{y \in \mathcal{C}} \langle y, \theta + \varepsilon Z \rangle$$

Follows a **perturbed model** with $Y \sim p_\theta(y)$, expectation $y_\varepsilon^*(\theta) = \mathbf{E}_{p_\theta}[Y]$.

Perturb and map Papandreou & Yuille (2011), FT Perturbed L Kalai & Vempala (2003)



Example. Over the unit simplex $\mathcal{C} = \Delta^d$ with Gumbel noise Z , $F(\theta) = \max_i \theta_i$.

$$F_\varepsilon(\theta) = \varepsilon \log \sum_{i \in [d]} e^{\frac{\theta_i}{\varepsilon}}, \quad p_\theta(e_i) \propto \exp(\langle \theta, e_i \rangle / \varepsilon), \quad [y_\varepsilon^*(\theta)]_i = \frac{e^{\frac{\theta_i}{\varepsilon}}}{\sum_j e^{\frac{\theta_j}{\varepsilon}}}$$

Why? and How?

Learning problems:

Features X_i , model output $\theta_w = g_w(X_i)$, prediction $y_{\text{pred}} = y_\varepsilon^*(\theta_w)$, loss L

$$F(w) = L(y_\varepsilon^*(\theta_w), y_i), \quad \text{gradients require } \partial_\theta y_\varepsilon^*(\theta_w).$$


Monte Carlo estimates. Perturbed maximizer and derivatives as expectations.

For $\theta \in \mathbf{R}^d$, $Z^{(1)}, \dots, Z^{(M)}$ i.i.d. copies

$$y^{(\ell)} = y^*(\theta + \varepsilon Z^{(\ell)})$$


Unbiased estimate of $y_\varepsilon^*(\theta)$ given by

$$\bar{y}_{\varepsilon, M}(\theta) = \frac{1}{M} \sum_{\ell=1}^M y^{(\ell)}.$$



Properties

Mirror maps: For \mathcal{C} with full interior, Z with smooth density μ , full support F_ε strictly convex, gradient Lipschitz. Ω strongly convex, Legendre type.



Differentiability. Functions are smooth in the inputs. For $\mu(z) \propto \exp(-\nu(z))$


$$y_\varepsilon^*(\theta) = \nabla_\theta F_\varepsilon(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[F(\theta + \varepsilon Z)\nabla_z \nu(Z)/\varepsilon],$$

$$\partial_\theta y_\varepsilon^*(\theta) = \nabla^2 F_\varepsilon(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)\nabla_z \nu(Z)^\top/\varepsilon].$$

Perturbed maximizer y_ε^* never locally constant in θ . Abernethy et al. (2014)

Properties

Mirror maps: For \mathcal{C} with full interior, Z with smooth density μ , full support F_ε strictly convex, gradient Lipschitz. Ω strongly convex, Legendre type.



Differentiability. Functions are smooth in the inputs. For $\mu(z) \propto \exp(-\nu(z))$


$$y_\varepsilon^*(\theta) = \nabla_\theta F_\varepsilon(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)] = \mathbf{E}[F(\theta + \varepsilon Z)\nabla_z \nu(Z)/\varepsilon],$$

$$\partial_\theta y_\varepsilon^*(\theta) = \nabla^2 F_\varepsilon(\theta) = \mathbf{E}[y^*(\theta + \varepsilon Z)\nabla_z \nu(Z)^\top/\varepsilon].$$

Perturbed maximizer y_ε^* never locally constant in θ . Abernethy et al. (2014)

Learning with perturbed optimizers

Machine learning pipeline: variable X , discrete label y , model outputs $\theta = g_w(X)$




Labels are solutions of optimization problems (one-hots, ranks, shortest paths)

Small modification of the model: end-to-end differentiable

Learning with perturbed optimizers

Machine learning pipeline: variable X , discrete label y , model outputs $\theta = g_w(X)$




Labels are solutions of optimization problems (one-hots, ranks, shortest paths)

Small modification of the model: end-to-end differentiable

Learning with perturbed optimizers

Machine learning pipeline: variable X , discrete label y , model outputs $\theta = g_w(X)$




Labels are solutions of optimization problems (one-hots, ranks, shortest paths)

Small modification of the model: end-to-end differentiable

Learning with perturbations and Fenchel-Young losses

Within the same framework, possible to virtually bypass the optimization block




Fenchel-Young losses Easier to implement, no Jacobian of y_ε^* . Blondel et al (20)

Population loss minimized at ground truth for perturbed generative model.

Learning with perturbations and Fenchel-Young losses

Motivated by model where $y_i = \operatorname{argmax}_{y \in \mathcal{C}} \langle g_{w_0}(X_i) + \varepsilon Z_i, y \rangle$



Stochastic gradients for empirical loss only require

$$\nabla_\theta L(\theta = g_w(X_i); y_i) = y_\varepsilon^*(\theta) - y_i = y_\varepsilon^*(g_w(X_i)) - y_i .$$

Simulated by a doubly stochastic scheme.

Computations

Monte Carlo estimates. Perturbed maximizer and derivatives as expectations.

For $\theta \in \mathbf{R}^d$, $Z^{(1)}, \dots, Z^{(M)}$ i.i.d. copies

$$y^{(\ell)} = y^*(\theta + \varepsilon Z^{(\ell)})$$

Unbiased estimate of $y_\varepsilon^*(\theta)$ given by


$$\bar{y}_{\varepsilon, M}(\theta) = \frac{1}{M} \sum_{\ell=1}^M y^{(\ell)}.$$

Supervised learning:

Features X_i , model output $\theta_w = g_w(X_i)$, prediction $y_{\text{pred}} = y_\varepsilon^*(\theta_w)$.

Stochastic gradient in w :

$$\nabla_w F_i(w) = \partial_w g_w(X_i) \cdot (y_\varepsilon^*(\theta) - Y_i)$$



Computations

Monte Carlo estimates. Perturbed maximizer and derivatives as expectations.

For $\theta \in \mathbf{R}^d$, $Z^{(1)}, \dots, Z^{(M)}$ i.i.d. copies

$$y^{(\ell)} = y^*(\theta + \varepsilon Z^{(\ell)})$$

Unbiased estimate of $y_\varepsilon^*(\theta)$ given by

$$\bar{y}_{\varepsilon, M}(\theta) = \frac{1}{M} \sum_{\ell=1}^M y^{(\ell)}.$$

Supervised learning:

Features X_i , model output $\theta_w = g_w(X_i)$, prediction $y_{\text{pred}} = y_\varepsilon^*(\theta_w)$.


Stochastic gradient in w (doubly stochastic scheme)

$$\nabla_w F_i(w) = \partial_w g_w(X_i) \cdot \left(\frac{1}{M} \sum_{\ell=1}^M y^*(\theta + \varepsilon Z^{(\ell)}) - Y_i \right).$$

Experiments

Learning to rank: Experiments on 4k instances of 100 vectors to rank.


Robustness to noise observed for some tolerated variance




Fenchel-Young loss is convex in w : linear model, possible theoretical analysis.

Experiments

Learning from shortest paths: From 10k examples of Warcraft 96×96 RGB images, representing 12×12 costs, and matrix of shortest paths. (Vlastelica et al. 19)





Train a CNN for 50 epochs, to learn costs recovery of optimal paths.



DEDAL

Deep Embedding and Differentiable ALignment

F. Llinares, Q. Berhet,
M. Blondel, O. Teboul,
J.P. Vert



- Deep embedding and alignment of protein sequences

Nature methods, 2023



F. Llinares



Q. Berthet



M. Blondel



O. Teboul




J-P. Vert

- Deep embedding and alignment of protein sequences

Nature methods, 2023


Protein alignment

- Learning character-wise embeddings of protein sequences.
- Using it to compute costs of an alignment problem (dynamic programming).



Protein alignment


- Learning character-wise embeddings of protein sequences.
- Using it to compute costs of an alignment problem (dynamic programming).




Differentiable alignment

- Given fixed substitution/insertion costs, local alignment problem.
- Smith-Waterman problem solved by DP, to align proteins.
- Non-differentiable solution, introducing perturbations

Alignment as a biologically-plausible sequence similarity measure



Source: [Dr. Vered Caspi](#)




Solved by dynamic program
(Smith-Waterman algorithm)

Differentiable alignment

- Given fixed substitution/insertion costs, local alignment problem.
- Smith-Waterman problem solved by DP, to align proteins.
- Non-differentiable solution, introducing perturbations




Differentiable Smith-Waterman alignments



Technique based on team's methodological work

Yields a differentiable version of the alignment algorithm



Confidential + Proprietary

P 8


Google Research

Differentiable alignment

- Given fixed substitution/insertion costs, local alignment problem.
- Smith-Waterman problem solved by DP, to align proteins.
- Non-differentiable solution, introducing perturbations




Differentiable Smith-Waterman alignments



Technique based on team's methodological work

Yields a differentiable version of the alignment algorithm

Google Research





Confidential + Proprietary

P 9

DEDAL: End-to-end learning to align


- From sequences to embeddings, costs, to perturbed alignments.
- Transformer architecture trained on databases of aligned proteins.






- **Self-supervised learning of audio representations from permutations with differentiable ranking**

A. Carr, Q. Berthet, M. Blondel, O. Teboul, N. Zeghidour
IEEE Signal Processing Letters, 2021



Rendering Le Lidec et al. (21)



Optimal Control Le Lidec et al. (22)

• Applications:

Digital pathology Thandiackal et al. (22), Patch selection Cordonnier et al. (21), Video Token Selection Wang et al. (22), . . .

• Algorithmic improvements: parallelized optimization Dubois-Taine et al. (22)



L. Stewart



F. Bach



F. Llinares




Q. Berthet


- **Differentiable Clustering with Perturbed Spanning Forests**
Preprint, 2023

Differentiable Clustering

- Transformer architecture trained on databases of aligned proteins.



- Semi-supervised clustering.
- Discovery of held-out classes.
- Presentation and poster** today.



Mahalo!

